

Sleep, Physical Activity, and Acid-Related Upper-GI Disease

Gastroesophageal Reflux Disease and Oesophagitis in the *All of Us* Research Program (Controlled Tier v9)

A Technical Report on Study Design, Data Engineering, and Statistical Modeling

Functional GI Disorders & Physiological Measures Project

July 18, 2026

Abstract

This report documents the design and implementation of a retrospective, cross-sectional analysis relating Fitbit-recorded *sleep* and *physical-activity* metrics (exposures) to two nested electronic-health-record (EHR) outcomes — broad gastroesophageal reflux disease (`has_gerd`) and its erosive *oesophagitis* subset (`has_esophagitis`) — among participants of the *All of Us* Research Program (Controlled Tier CDR v9, C2024Q3R9) who shared both wearable and EHR data. The study is organized as a 2×3 matrix: two outcomes crossed with three exposure sets (sleep, activity, and a combined activity+sleep set), giving six analysis configurations executed by one shared statistical engine. Six new R scripts implement the work without modifying any existing notebook: a dual-phenotype GERD outcome pull, a shared modeling-engine helper, three analysis files (sleep, activity, combined), and an optional proton-pump-inhibitor / H_2 -antagonist pull supporting a medication-treated “severe GERD” sensitivity analysis. This document details the phenotype definitions, cohort-eligibility logic for both wearable streams, exposure quartilization, covariate engineering (including the GERD-specific peptic-ulcer-disease and IBS covariates), the centralized descriptive / univariable / multivariable logistic-regression engine, mutually-adjusted combined models, assumption diagnostics, sensitivity analyses, and the ways the GERD pipeline both re-uses and departs from the project’s IBS and constipation templates. It is intended as a self-contained methodological reference for reproducing and critiquing the study.

Deliverables. GERD Outcome Data Upload R9.R, GERD Analysis Helpers R9.R, Sleep and GERD Analysis File R9.R, Activity and GERD Analysis File R9.R, Activity Sleep and GERD Analysis File R9.R, GERD Pharmacologics Data Upload.R.

Contents

1	Executive Summary	5
2	Introduction and Scientific Background	6
2.1	Acid-Related and Functional Upper Gastrointestinal Disorders	6
2.2	The Gut-Brain Axis and Behavioral-Physiological Exposures	6
2.3	Mechanisms Linking Poor Sleep and Low Physical Activity to Reflux Disease . .	7
2.4	Why Erosive Oesophagitis Is a Meaningful Stricter Phenotype	8
2.5	Rationale and the All of Us Research Opportunity	8
2.6	Research Question and Directional Hypothesis	9

3	Study Design Overview	9
3.1	Design at a Glance	10
3.2	The Six Analysis Configurations	10
3.3	Deliverable Files and Their Roles	11
3.4	Shared-Helper Architecture and Statistical Identity	11
4	Biological Mechanisms Linking Sleep and Activity to Reflux Disease	13
4.1	Lower-esophageal-sphincter tone and transient relaxations	13
4.2	Acid clearance and supine time	13
4.3	Intra-abdominal pressure and adiposity	13
4.4	Autonomic and vagal regulation	14
4.5	Visceral hypersensitivity and central pain processing	14
4.6	Circadian, inflammatory, and reverse-causal loops	14
5	Prior Evidence and the Knowledge Gap	14
6	Data Sources	14
6.1	The <i>All of Us</i> Controlled Tier, CDR v9	15
6.2	Wearable and EHR inputs	15
6.3	A key data fact	15
7	Outcome Ascertainment: GERD and Oesophagitis	16
7.1	Two nested phenotypes	16
7.2	Seed concepts and descendant expansion	16
7.3	The descendant-filter subtlety	16
7.4	Identifying the oesophagitis subset	17
7.5	Case-definition variants	17
8	Exposure Definitions	17
8.1	Sleep metrics	17
8.2	Activity metrics	17
8.3	Sleep×activity overlap sensitivity set	18
9	Combined Exposure Design	18
9.1	Intersection cohort	18
9.2	Covariate-timing anchor	18
9.3	Two model families	19
10	Cohort Construction and Eligibility	19
10.1	Eligibility gates	19
10.2	CONSORT accounting	19
11	Rationale for Eligibility and Cleaning Thresholds	19
12	Covariates	20
12.1	Temporal covariate rule	20
12.2	Covariate inventory	20
12.3	Medication, BMI, and GI covariate details	20
13	The Shared Modeling Engine and Descriptive Statistics	20
13.1	Engine functions	20
13.2	Descriptive statistics	21

14 Univariable and Multivariable Analysis	22
14.1 Univariable models	22
14.2 Multivariable models	22
14.3 Combined mutually-adjusted models	22
14.4 Assumption diagnostics	22
15 The Six Analysis Configurations in Detail	22
15.1 C1 — Sleep and broad GERD	23
15.2 C2 — Sleep and oesophagitis	23
15.3 C3 — Activity and broad GERD	23
15.4 C4 — Activity and oesophagitis	23
15.5 C5 — Combined exposures and broad GERD	23
15.6 C6 — Combined exposures and oesophagitis	23
16 The Logistic Model in Depth	24
16.1 Specification and estimation	24
16.2 Interpretation of the quartile contrasts	24
16.3 Variable selection and the backward model	24
16.4 Discrimination and multicollinearity	24
17 Model Assumptions and Remediation	24
18 Data Engineering, Quality Control, and Statistical Considerations	25
18.1 One-row-per-participant guarantee	25
18.2 Missing data	25
18.3 Multiple comparisons	25
18.4 Confounding structure	25
18.5 Sample size and precision	26
19 Sensitivity Analyses	26
19.1 The pharmacologics pull	26
20 Relationship to the IBS and Constipation Templates	26
21 Reproducibility and Execution	26
21.1 Execution order	26
21.2 Software environment	27
22 Anticipated Results and Interpretation Framework	27
23 Threats to Validity	28
23.1 Selection bias	28
23.2 Information (measurement) bias	28
23.3 Confounding	28
23.4 Reverse causation	28
23.5 External validity	29
24 Limitations	29
25 Future Directions and Extensions	29
26 Integration Across the FGID Portfolio	30
27 Conclusion	30

A	Appendix A: Output Inventory	31
B	Appendix B: Concept Reference	31
C	Appendix C: Selected Code Excerpts	32
D	Appendix D: Glossary of Key Derived Variables	32
E	Appendix E: Results-Template Tables	33
E.1	Adjusted odds ratios — sleep exposures (C1 broad GERD; C2 oesophagitis) . . .	33
E.2	Adjusted odds ratios — activity exposures (C3 broad GERD; C4 oesophagitis) .	33
E.3	Mutually-adjusted combined pairs (C5 broad GERD; C6 oesophagitis)	33
E.4	Descriptive-table shell (per outcome)	34
F	Appendix F: Per-File Walkthrough	34
F.1	GERD Outcome Data Upload R9.R	34
F.2	GERD Analysis Helpers R9.R	34
F.3	Sleep and GERD Analysis File R9.R	34
F.4	Activity and GERD Analysis File R9.R	35
F.5	Activity Sleep and GERD Analysis File R9.R	35
F.6	GERD Pharmacologics Data Upload.R	35
G	Appendix G: Reporting and Reproducibility Checklist	36

1 Executive Summary

This report specifies a retrospective, cross-sectional analysis conducted in the All of Us Research Program (Controlled Tier Curated Data Repository v9, release C2024Q3R9) that relates objectively measured, Fitbit-recorded physiological exposures – sleep and physical activity – to acid-related upper gastrointestinal disease ascertained from electronic health records (EHR). The study is deliberately built around *two* binary outcome phenotypes and *three* exposure sets, yielding a 2×3 grid of six primary analysis configurations that share a single statistical engine.

The two outcomes are nested EHR phenotypes, each requiring at least two condition records (≥ 2). The broad outcome, `has_gerd`, captures gastroesophageal reflux disease from all rows returned by a dedicated GERD condition pull. The stricter outcome, `has_esophagitis`, is the erosive subset – GERD records whose standard concept name contains “esophagitis” or “oesophagitis” – and is seeded conceptually by concept 4223293 (“GERD with esophagitis”). Because `has_esophagitis` is a strict subset nested within the same GERD-record pull, it represents a rarer, more specific phenotype of objective mucosal injury that is analyzed in parallel with the broad, largely symptom-coded `has_gerd`.

The three exposure sets are each analyzed against *both* outcomes. (A) Sleep exposures are cohort quartiles (with Q1 as the reference level) of average nightly minutes asleep, minutes in bed, minutes awake, minutes restless, deep, light, and REM sleep, and sleep efficiency (asleep divided by in-bed). (B) Activity exposures are cohort quartiles of average daily steps, lightly, fairly, very, and total active minutes, sedentary minutes, and a corrected maximum daily heart rate. (C) The combined set operates on the intersection cohort (participants with valid sleep *and* valid activity streams) and analyzes each of the roughly fifteen exposures separately, plus mutually-adjusted models that carry a sleep quartile and an activity quartile together in the same equation (for example, minutes-asleep quartile jointly with steps quartile).

Exposure set	Outcome <code>has_gerd</code>	Outcome <code>has_esophagitis</code>
(A) Sleep quartiles	Config 1	Config 2
(B) Activity quartiles	Config 3	Config 4
(C) Combined (+ mutually adjusted)	Config 5	Config 6

Table 1: The six analysis configurations formed by crossing three exposure sets with two nested outcome phenotypes. Every configuration is fit with the identical modeling engine.

To guarantee that all six configurations use identical statistics, definitions, and reporting, the design centralizes methodology in a shared engine file, “GERD Analysis Helpers R9.R”. That helper defines the outcome objects and factor labels (`GERD_OUTCOMES` with “No GERD”/“GERD” and “No oesophagitis”/“Oesophagitis”), the exposure lists (`SLEEP_EXPOSURES`, `ACTIVITY_EXPOSURES`), the adjustment set (`GERD_ADJ_COVARS`, which adds `has_pud` and `has_ibs` relative to the antecedent IBS study and uses the Charlson Comorbidity Index categorically by default), the master label lookup (`GERD_LABELS`), and the reusable functions `prep_modeling_df()`, `descriptive_table()`, `univariate_table()`, `run_models_for_outcome()`, and `run_diagnostics()`. A single dispatcher, `analyze_all_outcomes()`, loops over both outcomes so that each analysis file merely builds its data frame, sources the helper, and calls the dispatcher.

The deliverables are all new R files; no existing notebook is edited:

- “GERD Outcome Data Upload R9.R” – the dual-phenotype outcome pull (seeds concepts 318800 and 4223293, expanding to standard descendants), producing `R9_gerd_outcome.rds`.
- “GERD Analysis Helpers R9.R” – the shared modeling engine described above.
- “Sleep and GERD Analysis File R9.R” – sleep exposures against both outcomes.

- “Activity and GERD Analysis File R9.R” – activity exposures against both outcomes.
- “Activity Sleep and GERD Analysis File R9.R” – the combined cohort, both outcomes, plus mutually-adjusted models.
- “GERD Pharmacologics Data Upload.R” – an optional PPI/H2RA prescription pull supporting a medication-treated (“severe”) GERD sensitivity analysis, producing `R9_gerd_drug_df.rds`.

Primary analytic outputs are `R9_final_analysis_sleep_gerd_df.rds`, `R9_final_analysis_activity_gerd_df.rds`, and `R9_final_analysis_activity_sleep_gerd_df.rds`, together with the reusable status objects `R9_gerd_status.rds` and `R9_esophagitis_status.rds`. All estimates are adjusted odds ratios from logistic regression, quartile exposures are compared against Q1, and multiplicity across the many exposure–outcome contrasts is controlled with Benjamini–Hochberg false-discovery-rate (FDR) q -values. Because the design is cross-sectional, the primary estimates are associations; a pre-specified post-Fitbit temporal outcome definition serves as the strongest guard against reverse causation.

2 Introduction and Scientific Background

2.1 Acid-Related and Functional Upper Gastrointestinal Disorders

Gastroesophageal reflux disease (GERD) is among the most prevalent chronic gastrointestinal conditions in adults, defined clinically as troublesome symptoms or mucosal complications arising when gastric contents reflux into the oesophagus. It sits at the intersection of two overlapping disease constructs. On one side are the acid-related structural disorders, in which excess oesophageal acid exposure produces objective epithelial injury – erosions, ulceration, stricture, and Barrett metaplasia. On the other side are the functional and symptom-defined presentations, in which patients experience heartburn, regurgitation, and non-cardiac chest pain in the absence of, or out of proportion to, visible mucosal damage. This latter category shares mechanistic and epidemiologic features with the functional gastrointestinal disorders (now termed disorders of gut–brain interaction), including a prominent role for visceral hypersensitivity, altered central pain processing, and psychosocial and behavioral modulation of symptom burden.

This dual character motivates the study’s two nested phenotypes. Broad GERD (`has_gerd`) is dominated by symptom-coded and clinically diagnosed reflux and therefore behaves partly as a disorder of gut–brain interaction, whereas erosive oesophagitis (`has_esophagitis`) isolates the subset with documented, objective mucosal injury. Studying both in parallel lets us ask whether physiological exposures act on subjective symptom reporting, on measurable tissue damage, or on both, and whether the strength of association differs between the sensory/perceptual and the structural manifestations of reflux.

2.2 The Gut-Brain Axis and Behavioral-Physiological Exposures

The gut–brain axis is the bidirectional signaling network linking the central nervous system, the autonomic nervous system, the enteric nervous system, and the gut. Through this axis, higher-order behavioral states – sleep, physical activity, arousal, and stress – exert continuous influence over gastrointestinal motility, secretion, mucosal blood flow, and afferent pain signaling, while visceral signals in turn shape sleep architecture and perceived exertion. Sleep and physical activity are therefore not incidental lifestyle covariates but plausible upstream physiological modulators of reflux pathophysiology.

Historically, these exposures have been measured by self-report, which is subject to recall bias, social desirability, and poor correspondence with objectively recorded behavior. Consumer wearable devices such as Fitbit allow longitudinal, objective, per-night and per-day quantification

of sleep quantity and quality and of ambulatory activity intensity and volume. This study exploits that measurement richness: rather than a single crude exposure, it decomposes sleep into duration, continuity, and stage architecture (deep, light, REM, awake, restless, efficiency) and activity into volume (steps), intensity strata (lightly, fairly, very active minutes), sedentary time, and a fitness/exertion proxy (corrected maximum daily heart rate). Quartile categorization (Q1 reference) preserves monotone dose–response signal while remaining robust to the skew and device-specific scaling of these metrics.

2.3 Mechanisms Linking Poor Sleep and Low Physical Activity to Reflux Disease

Several partially overlapping mechanisms make both insufficient or fragmented sleep and low physical activity biologically plausible contributors to GERD and, more specifically, to erosive oesophagitis.

Supine nocturnal acid exposure. Sleep is spent predominantly recumbent, and the supine posture removes the gravitational assistance that clears refluxate during waking hours. Nocturnal reflux episodes tend to be longer and to reach more proximal segments of the oesophagus, prolonging contact time between acid and the mucosa. Individuals with shorter, more fragmented, or more inefficient sleep may accumulate greater cumulative supine acid-contact time, and disturbed sleep architecture is itself associated with a higher frequency of nocturnal reflux events. Because prolonged acid contact is the proximate driver of epithelial erosion, this pathway predicts a particularly strong link to the erosive `has_esophagitis` phenotype.

Delayed acid clearance. Oesophageal acid clearance depends on primary peristalsis triggered by swallowing and on neutralization by swallowed, bicarbonate-rich saliva. Both swallowing frequency and salivary flow fall markedly during sleep, and this suppression is exaggerated by sleep deprivation and fragmentation. The result is slower restoration of normal intraluminal pH after a reflux event, again lengthening mucosal acid exposure. Reduced daytime physical activity may compound impaired motility indirectly through its downstream metabolic and autonomic effects.

Visceral hypersensitivity. Experimental sleep deprivation lowers the oesophageal pain threshold, producing hyperalgesia to acid and mechanical distension. This heightened afferent sensitivity increases symptom reporting for a given quantum of acid exposure and is a core feature of the gut–brain-interaction component of reflux. This pathway predicts that poor sleep will associate with symptom-defined `has_gerd` even where objective injury is modest, and it is a mechanism through which sleep can influence perceived disease independently of tissue damage.

Autonomic and vagal tone. Both habitual physical activity and healthy sleep shape autonomic balance, and vagal tone modulates lower oesophageal sphincter (LES) pressure and the frequency of transient LES relaxations (TLESRs), the dominant mechanism of reflux events. Disrupted sleep and physical deconditioning shift autonomic balance in directions associated with impaired sphincter function and altered gastric accommodation and emptying. The corrected maximum daily heart rate is included in part as a proxy that carries autonomic and cardiorespiratory-fitness information alongside the volume and intensity activity metrics.

Adiposity and mechanical load. Low physical activity is a determinant of positive energy balance, weight gain, and especially central (visceral) adiposity. Increased intra-abdominal pressure raises the gastro-oesophageal pressure gradient, promotes hiatal hernia formation, and disrupts the anti-reflux barrier – a well-established mechanical route from sedentary behavior to

reflux and to erosive disease. Adiposity is therefore both a mediator on the causal path from inactivity to GERD and a confounder to be handled analytically; the design measures median BMI (from measurement concept 3038553) and adjusts for it, recognizing that this partially blocks a genuine causal pathway and may attenuate activity–outcome estimates.

Reverse causation. Critically, the sleep–reflux relationship is bidirectional: nocturnal acid reflux itself provokes arousals and fragments sleep, so poor sleep can be a consequence rather than a cause of GERD. Likewise, symptomatic or severe reflux can discourage vigorous exertion, and post-prandial vigorous activity can transiently provoke reflux, so the activity–reflux relationship is not unidirectional either. In a cross-sectional design these bidirectional effects cannot be fully disentangled, which is why the analysis pre-specifies a post-Fitbit temporal outcome definition (first GERD code recorded ≥ 180 days after the first Fitbit date, then ≥ 2 codes) as its principal guard against reverse causation, complemented by distinct-date, overlap-set, and medication-treated sensitivity analyses.

2.4 Why Erosive Oesophagitis Is a Meaningful Stricter Phenotype

Broad GERD ascertained from EHR condition codes is heterogeneous: it aggregates objective erosive disease with a large fraction of symptom-driven, non-erosive, and empirically treated reflux, and it is susceptible to differential coding by healthcare utilization and to misclassification of functional heartburn as GERD. Erosive oesophagitis (`has_esophagitis`), by contrast, requires documentation consistent with visible mucosal injury and thus anchors the analysis to a more objective, endoscopically grounded endpoint. Several considerations make it a valuable second phenotype rather than a redundant one.

First, the two phenotypes plausibly load on different mechanisms: symptom-defined GERD is more strongly shaped by visceral hypersensitivity and central pain processing (sensitive to sleep-related hyperalgesia), whereas erosive disease is more directly a function of cumulative acid-contact time and barrier failure (sensitive to supine exposure, delayed clearance, and adiposity). Contrasting the exposure associations across the two outcomes is therefore mechanistically informative. Second, erosive oesophagitis is less vulnerable to the diagnostic-labeling noise that inflates symptom-coded GERD, so a concordant signal across both phenotypes strengthens causal interpretation, while a signal confined to symptom-defined GERD points toward a perceptual rather than a structural pathway. Third, because `has_esophagitis` is a strict subset nested within the same GERD-record pull, it is by construction rarer than broad GERD; this reduces statistical power and can widen confidence intervals, which the analysis accommodates through complete-case reporting, FDR control, and diagnostic checks for separation and unstable standard errors. Reporting both phenotypes with the identical modeling engine makes the comparison fair and transparent.

2.5 Rationale and the All of Us Research Opportunity

Prior evidence linking sleep and activity to reflux is dominated by cross-sectional self-report, small clinical samples, and endpoints that rarely separate symptom-defined from erosive disease. The All of Us Research Program provides a distinctive opportunity to address these gaps at scale: it links longitudinal, device-recorded Fitbit sleep and activity data to a rich, standardized (OMOP-mapped) EHR and to detailed survey and physical-measurement data across a large, deliberately diverse participant population. This enables objectively measured exposures, EHR-based outcome phenotyping with an explicit ≥ 2 -record rule, adjustment for an extensive covariate set (demographics, socioeconomic factors, alcohol, smoking, BMI, Charlson Comorbidity Index, and a panel of comorbidities and relevant medication-use flags), and reflux-specific confounders including peptic ulcer disease (`has_pud`) and irritable bowel syndrome (`has_ibs`), the latter reflecting the known overlap of reflux with other disorders of gut–brain interaction.

The design mirrors the project’s antecedent IBS sleep study, constipation activity study, and constipation outcome upload, but adapts them for reflux: it swaps the IBS outcome for the GERD outcomes and adds the erosive subset, sources outcomes from a dedicated GERD pull (necessary because `R9_condition_df.rds` does not contain GERD), removes all IBS-subtype and multinomial logic, adds `has_pud` and `has_ibs` as covariates, harmonizes `R9_*` file naming, extends the work to activity and combined analyses, and centralizes all statistics in the shared helper so that every configuration is estimated identically.

2.6 Research Question and Directional Hypothesis

The central research question is whether objectively measured sleep and physical activity are associated with the odds of broad GERD and of erosive oesophagitis in Fitbit-sharing All of Us participants, and whether any such associations persist after adjustment for demographic, socioeconomic, anthropometric, comorbidity, and medication confounders. The directional hypothesis is that greater sleep quantity and quality (for example, higher minutes asleep, higher sleep efficiency, and more consolidated sleep) and greater physical activity (for example, higher daily steps and more active minutes, with correspondingly lower sedentary time) associate with *lower* odds of both GERD and oesophagitis – that is, upper quartiles (Q2–Q4) will carry adjusted odds ratios below unity relative to Q1. We anticipate that these associations will attenuate after covariate adjustment – particularly for activity exposures, where BMI lies partly on the causal pathway – but that they will persist rather than disappear. Given the mechanistic profile described above, we expect the acid-exposure-dependent pathways to register on the erosive `has_esophagitis` phenotype, while sleep-related hyperalgesia may leave a signal on symptom-defined `has_gerd` even where mucosal injury is limited. Because the design is cross-sectional, all primary estimates are interpreted as associations; the post-Fitbit temporal outcome definition provides the strongest available guard against the reverse-causation pathway in which nocturnal reflux itself fragments sleep.

3 Study Design Overview

This study is a retrospective, cross-sectional analysis conducted in the All of Us Research Program using the Controlled Tier Curated Data Repository (CDR) v9, release C2024Q3R9. It relates wearable-recorded physiological exposures, derived from Fitbit devices, to acid-related upper gastrointestinal disease captured in the electronic health record (EHR). The analysis estimates associations between habitual sleep and physical-activity patterns and the odds of two nested EHR-defined outcomes: broad gastroesophageal reflux disease (`has_gerd`) and the more specific erosive subset (`has_esophagitis`). Because the design is cross-sectional, the primary estimands are associations rather than causal effects; a post-Fitbit temporal outcome definition is carried as the principal guard against reverse causation, since nocturnal reflux is itself known to fragment sleep.

The design is organized around a factorial matrix of two outcomes crossed with three exposure sets, yielding six analysis configurations that are executed by a common statistical engine. The two outcomes are binary EHR phenotypes requiring ≥ 2 qualifying condition records. The three exposure sets are (A) sleep metrics, (B) activity metrics, and (C) a combined activity-and-sleep set evaluated on the intersection cohort, including mutually adjusted models that carry one sleep and one activity exposure simultaneously. Every exposure is operationalized as a cohort quartile factor with the first quartile (Q1) as the reference level, and every configuration is fit with the identical multivariable logistic-regression specification described in Section 3.4 of the Methods (the shared modeling engine). The remainder of this section presents a design-at-a-glance summary, enumerates the six configurations and the deliverable code, and explains the don’t-repeat-yourself (DRY) architecture that guarantees statistical identity across all six.

3.1 Design at a Glance

Table 2 summarizes the core design parameters. The study population is the set of Fitbit-sharing All of Us participants who are age ≥ 18 at their first relevant Fitbit date, who share EHR data (condition or measurement records), and who meet the stream-specific wear-time and valid-record thresholds detailed in the Data and Cohort sections. Exposures are cohort-internal quartiles so that each configuration compares the upper three quartiles of an exposure against its own Q1 reference within the analytic cohort for that exposure set.

Design element	Specification
Data source	All of Us Research Program, Controlled Tier CDR v9 (release C2024Q3R9)
Study design	Retrospective, cross-sectional, EHR outcomes with wearable-derived exposures
Population	Fitbit-sharers; age ≥ 18 at first Fitbit date; sharing EHR (condition or measurement)
Exposures	Cohort quartiles (Q1 reference) of Fitbit sleep and activity metrics; combined set adds mutually adjusted sleep + activity models
Outcomes	has_gerd (broad GERD, ≥ 2 records) and has_esophagitis (erosive subset, ≥ 2 records, concept-name match), a nested phenotype
Primary model	Multivariable logistic regression; $\exp(\beta_j)$ is the adjusted odds ratio for quartile j versus Q1
Covariate anchor	Pre-Fitbit window; combined cohort anchored to the earliest Fitbit date across sleep and activity streams (pmin)
Covariates	Demographics, socioeconomic status, median BMI, Charlson Comorbidity Index, comorbidities incl. has_pud and has_ibs , and medication-use flags
Missing data	Binary EHR indicators coalesced NA $\rightarrow$ FALSE; survey/measurement missingness retained as an explicit Missing level
Multiplicity	Benjamini–Hochberg false discovery rate (FDR) q -values
Configurations	Six: 2 outcomes $\times$ 3 exposure sets, all fit by one shared engine
Sensitivity analyses	Distinct-date GERD; post-Fitbit temporal GERD; sleep $\times$ activity overlap set; optional medication-treated (severe) GERD

Table 2: Design at a glance for the GERD and oesophagitis wearable-exposure analysis.

3.2 The Six Analysis Configurations

The analytic space is the 2×3 crossing of outcomes and exposure sets shown in Table 3. Rows correspond to the two nested outcomes and columns correspond to the three exposure sets. Each of the three exposure sets is implemented by a single analysis file that builds one modeling data frame and then loops over both outcomes via the shared driver `analyze_all_outcomes()`; consequently the two configurations in a given column share the same analysis file and the same primary output data frame, differing only in which outcome column the engine stratifies and models. This is why the file and primary output are common within a column: the outcome loop is internal to a single script, not split across scripts.

The combined column (C5, C6) is defined on the intersection cohort of participants who satisfy both the sleep-validity threshold (≥ 180 valid nights) and the activity-validity threshold (≥ 30 valid days). In that column each of the roughly fifteen exposures is analyzed on its own and, in addition, mutually adjusted models are fit that include one sleep-metric quartile and one

Outcome	(A) Sleep exposures	(B) Activity exposures	(C) Combined (activity + sleep)
has_gerd (broad GERD)	C1. File Sleep and GERD Analysis File R9.R; output R9_final_analysis_sleep_gerd_df.rds	C3. File Activity and GERD Analysis File R9.R; output R9_final_analysis_activity_gerd_df.rds	C5. File Activity Sleep and GERD Analysis File R9.R; output R9_final_analysis_activity_sleep_gerd_df.rds
has_esophagitis (erosive subset)	C2. Same file and primary output as C1 (second outcome in the loop)	C4. Same file and primary output as C3 (second outcome in the loop)	C6. Same file and primary output as C5 (second outcome in the loop)

Table 3: The six analysis configurations (rows: outcomes; columns: exposure sets), naming the analysis file and primary output data frame for each. Both configurations in a column are produced by the same file because `analyze_all_outcomes()` loops both outcomes over one modeling data frame.

activity-metric quartile together (for example, minutes-asleep quartile jointly with steps quartile). Covariate timing for the combined cohort is anchored to each participant’s earliest Fitbit date across the two streams, computed as the element-wise minimum (`pmin`) of the first sleep date and the first activity date, so that all pre-Fitbit covariate windows are defined consistently regardless of which stream began first.

3.3 Deliverable Files and Their Roles

The project is delivered as six new R scripts, listed in Table 4. Two are data-upload scripts that materialize outcome and pharmacologic data via BigQuery / Cohort Builder pulls; one is the shared modeling engine; and three are the analysis files that build modeling data frames and invoke the engine. The primary modeling outputs are the three `R9_final_analysis_*_gerd_df.rds` data frames plus the reusable outcome-status objects `R9_gerd_status.rds` and `R9_esophagitis_status.rds`.

A dedicated outcome pull is required because the existing comorbidity object `R9_condition_df.rds` does not contain GERD condition records. The outcome upload therefore seeds the two GERD concepts and expands to all standard descendants using the `cb_criteria` path-match pattern. Because that pull is already GERD-only, `has_gerd` is computed over all pulled rows rather than by re-filtering on only the two seed concept identifiers; re-filtering on the seeds alone would incorrectly drop descendant-coded cases. `has_esophagitis` is then identified within the same pull by a concept-name match on “esophagitis” / “oesophagitis” (seeded by concept 4223293), making it a strict subset phenotype nested within the broad GERD pull and, as the erosive definition, rarer than broad GERD.

3.4 Shared-Helper Architecture and Statistical Identity

All statistical logic is centralized in a single helper file, `GERD Analysis Helpers R9.R`, which each analysis file loads with `source()` before calling the common driver. This don’t-repeat-yourself (DRY) architecture exists specifically to guarantee that all six configurations are estimated with byte-for-byte identical statistics: the only quantities allowed to differ across configurations are the input modeling data frame and the exposure list handed to the engine.

The helper file is the single canonical definition of every modeling choice. It defines:

- `GERD_OUTCOMES` – the two outcomes together with their factor labels (“No GERD”/“GERD” and “No oesophagitis”/“Oesophagitis”), so outcome coding is fixed once.

Deliverable file	Role
GERD Outcome Data Upload R9.R	Dual-phenotype outcome pull; seeds concepts 318800 (GERD) and 4223293 (GERD with esophagitis) and expands to all standard descendants, restricted to Fitbit-sharers, producing <code>R9_gerd_outcome.rds</code> , <code>R9_gerd_status.rds</code> , and <code>R9_esophagitis_status.rds</code>
GERD Analysis Helpers R9.R	Shared modeling engine (DRY): defines outcomes, exposure lists, adjustment covariates, label lookups, and all descriptive, univariate, model-fitting, and diagnostic functions
Sleep and GERD Analysis File R9.R	Builds the sleep modeling data frame and runs both outcomes; primary output <code>R9_final_analysis_sleep_gerd_df.rds</code> (configurations C1, C2)
Activity and GERD Analysis File R9.R	Builds the activity modeling data frame and runs both outcomes; primary output <code>R9_final_analysis_activity_gerd_df.rds</code> (configurations C3, C4)
Activity Sleep and GERD Analysis File R9.R	Builds the combined intersection-cohort data frame, runs both outcomes and mutually adjusted models; primary output <code>R9_final_analysis_activity_sleep_gerd_df.rds</code> (configurations C5, C6)
GERD Pharmacologics Data Upload.R	Optional PPI (ancestor 21600095) and H2RA (ancestor 21600096) prescription pull producing <code>R9_gerd_drug_df.rds</code> , used only by the medication-treated (severe) GERD sensitivity analysis and guarded by file existence

Table 4: Deliverable code files and their roles. All six files are new; no existing All of Us notebook is edited.

- `SLEEP_EXPOSURES` and `ACTIVITY_EXPOSURES` – the canonical exposure inventories consumed by the sleep, activity, and combined files.
- `GERD_ADJ_COVARS` – the shared adjustment set (which adds `has_pud` and `has_ibs` relative to the sibling IBS study and uses the categorical Charlson Comorbidity Index by default), so every model adjusts for the same covariates.
- `GERD_LABELS` – the master label lookup used in all output tables.
- `prep_modeling_df()` – builds the collapsed factor covariates and relevels each quartile exposure to the $Q1 \rightarrow Q4$ ordering with $Q1$ as reference.
- `descriptive_table()`, `univariate_table()`, `run_models_for_outcome()`, `run_diagnostics()`, and `analyze_all_outcomes()` – the descriptive, unadjusted, adjusted-model, diagnostic, and both-outcome driver routines.

Because these objects and functions live in exactly one place, every configuration inherits the same estimator (logistic regression), the same covariate specification, the same reference-level convention, the same complete-case handling, the same full-model-plus-backward-AIC pairing (via `MASS::stepAIC` with the exposure forced to remain), the same area-under-the-curve computation via `pROC`, the same variance-inflation-factor computation via `car`, the same descriptive stratified tables with chi-square tests and FDR q -values, and the same diagnostic battery (independence, multicollinearity, separation and large standard errors, Cook’s distance, and Box–Tidwell linearity for BMI). The driver `analyze_all_outcomes()` loops both outcomes so that a single call in each analysis file produces both configurations in that exposure column under identical code paths.

This centralization delivers three concrete guarantees. First, statistical identity: any two configurations that should be comparable are, by construction, computed by the same functions with the same defaults, eliminating silent divergence between the sleep, activity, and combined analyses. Second, single-point maintainability: a change to the covariate set, a label, or a diagnostic propagates to all six configurations at once, so the report cannot fall out of internal consistency. Third, auditability: the exact model that produced any table can be traced to one helper definition rather than to per-file copies that might drift. The design mirrors the project's IBS sleep study, constipation activity study, and constipation outcome upload, but removes all IBS-subtype logic (subtype flags and multinomial models), swaps the IBS covariate for the GERD outcomes plus the added `has_esophagitis`, adds `has_pud` and `has_ibs` as covariates, standardizes `R9_naming`, extends the analysis to activity and combined exposures, and centralizes the modeling engine in the shared helper.

Finally, all deliverables are new files. No existing All of Us notebook or script is modified: the outcome and pharmacologic pulls, the shared helper, and the three analysis files are added alongside the existing project code, and prior objects such as `R9_condition_df.rds` are read but never rewritten. This isolation keeps the GERD analysis reproducible and non-destructive with respect to the established IBS and constipation pipelines from which it inherits its templates.

4 Biological Mechanisms Linking Sleep and Activity to Reflux Disease

Because the study relates two behavioral-physiological exposures to two reflux phenotypes, it is worth setting out, in one place, the candidate mechanisms that make each exposure–outcome link plausible. These mechanisms also motivate the covariate set (each major confounder sits on one of these pathways) and the directional predictions of Section 19.

4.1 Lower-esophageal-sphincter tone and transient relaxations

Reflux events are driven largely by transient lower-esophageal-sphincter relaxations. Sleep deprivation, sympathetic arousal, and the substances associated with poor sleep and sedentary lifestyles (nicotine, alcohol, late large meals) each lower resting sphincter tone or increase the frequency of transient relaxations, raising the probability that gastric contents cross into the oesophagus. Adjusting for smoking and alcohol partially isolates the sleep/activity signal from these co-travelling behaviors.

4.2 Acid clearance and supine time

Once refluxate reaches the oesophagus, clearance depends on primary peristalsis (swallowing) and salivary bicarbonate, both of which are suppressed during sleep. Fragmented or short sleep can extend supine, low-clearance time and prolong mucosal acid contact, a pathway that plausibly matters more for the erosive oesophagitis phenotype (`has_esophagitis`) than for symptom-coded GERD, because sustained contact time is what produces mucosal injury.

4.3 Intra-abdominal pressure and adiposity

Higher body-mass index raises intra-abdominal pressure, promotes hiatal herniation, and increases the gastro-oesophageal pressure gradient favoring reflux. Adiposity is also associated with shorter, more fragmented sleep and with lower physical activity, making it a central confounder — which is why median BMI is included in every adjusted model and why any residual sleep/activity association is interpreted as over-and-above the adiposity pathway.

4.4 Autonomic and vagal regulation

Sleep and exercise both modulate autonomic balance. Vagal tone governs gastric motility and lower-esophageal-sphincter function; poor sleep shifts autonomic balance sympathetically, and habitual activity raises vagal tone. The corrected maximum-heart-rate exposure in the activity set is, in part, a proxy for this autonomic/fitness axis.

4.5 Visceral hypersensitivity and central pain processing

Sleep deprivation lowers visceral pain thresholds, so the same acid exposure can produce more symptoms. This pathway acts preferentially on the symptom-defined broad GERD phenotype rather than on objective mucosal injury, and it is the principal reason the two outcomes are analyzed in parallel: an exposure that acts through perception should show a stronger association with `has_gerd` than with `has_esophagitis`, whereas one acting through acid contact time should show the reverse or an equal association.

4.6 Circadian, inflammatory, and reverse-causal loops

Circadian disruption alters gastric acid secretion and melatonin (which has mucosal-protective effects), and low-grade systemic inflammation associated with poor sleep and inactivity may sensitize the oesophageal mucosa. Finally, the loop is bidirectional: nocturnal reflux itself awakens patients and fragments sleep, which is precisely why the post-Fitbit temporal outcome is carried as the primary reverse-causation guard.

5 Prior Evidence and the Knowledge Gap

Associations between lifestyle factors and reflux have been studied for decades, but the evidence base has three recurring limitations that this study is designed to address. First, most prior work measures *sleep and activity by self-report*, which is subject to recall and social-desirability bias and correlates only modestly with objective wearable measurement; here both exposures are device-recorded at nightly and daily resolution over months. Second, reflux outcomes are often ascertained by *symptom questionnaire in modest single-center samples*, limiting both objectivity and power; here the outcome is EHR-defined across a large, national, demographically broad cohort, and the erosive oesophagitis subset provides an objective, mucosal-injury endpoint alongside the symptom-coded broad phenotype. Third, few datasets *link high-resolution wearable streams to longitudinal EHR* at scale; the *All of Us* Research Program is purpose-built for exactly this linkage.

The specific gap this study fills is therefore the near-absence of large-scale, objectively-measured sleep-and-activity analyses of GERD — and, in particular, the lack of any analysis that separates the symptom-coded and erosive phenotypes while adjusting for an unusually complete covariate set including adiposity, psychiatric comorbidity, peptic ulcer disease, and IBS. By running the identical engine over sleep, activity, and combined exposures and over both outcomes, the design also permits an internal triangulation — across exposure domains and across outcome phenotypes — that single-exposure, single-outcome studies cannot provide. No numeric findings are asserted here; the contribution of this report is the specification of that analysis.

6 Data Sources

6.1 The *All of Us* Controlled Tier, CDR v9

All queries and derived data use the Controlled Tier Curated Data Repository, version 9 (C2024Q3R9). The outcome and drug pulls pin the workspace CDR explicitly with `Sys.setenv(WORKSPACE_CDR = "fc-aou-cdr-prod-ct.C2024Q3R9")` so that concept expansions resolve against the intended release. All person-level data remain inside the secure *All of Us* workspace; only code is version-controlled, and this report contains no individual-level data.

6.2 Wearable and EHR inputs

The analysis draws on three Fitbit streams and four EHR/survey domains. The sleep stream provides nightly stage minutes; the intraday-steps stream provides per-day step totals and wear hours and defines valid activity days; the daily activity stream provides activity-zone minutes; and a pre-computed corrected maximum-heart-rate summary supplies the cardiovascular exposure. The condition, measurement, drug, survey, and person tables supply the outcome (indirectly), covariates, and demographics. Table 5 lists every input.

Input .rds	Contents / role
R9_gerd_outcome.rds	New. All GERD condition records for Fitbit-sharers; source of both outcomes.
R9_fitbit_sleep_daily_summary_df.rds	Nightly Fitbit sleep summaries (sleep exposure source).
R9_fitbit_activity_df.rds	Daily activity-zone minutes (sedentary, lightly/fairly/very active).
R9_fitbit_intraday_steps_df.rds	Intraday steps and wear hours; defines valid activity days.
R9_max_hr_minute_all_days_df.rds	Corrected maximum daily heart rate (optional exposure).
R9_person_df.rds	Demographics (date of birth, sex at birth, race, ethnicity, gender).
R9_survey_df.rds	Survey answers (smoking, alcohol, education, income).
R9_measurement_df.rds	Measurements, including body mass index (concept 3038553).
R9_drug_df.rds	Drug exposures for cardiac/psychiatric medication covariates.
R9_condition_df.rds	Comorbidity conditions (and IBS); <i>does not</i> contain GERD.
R9_pud_status.rds	Pre-computed peptic-ulcer-disease covariate (<code>has_pud</code>).
cci_scored_sleep.rds / R9_cci_scored.rds	Charlson Comorbidity Index scores.
R9_gerd_drug_df.rds	New, optional. PPI/ H_2 RA exposures (treated-GERD analysis).

Table 5: Datasets consumed by the GERD analyses. Only `R9_gerd_outcome.rds` (and optionally `R9_gerd_drug_df.rds`) must be created; everything else is reused as-is.

6.3 A key data fact

`R9_condition_df.rds` was pulled with concept sets for depression, anxiety, diabetes, hypertension, COPD, obstructive sleep apnea, heart failure, myocardial infarction, stroke, and IBS (concept

75576)—but *not* GERD. GERD must therefore be ascertained through a dedicated Cohort Builder export; it cannot be recovered from the existing condition table. Peptic ulcer disease is likewise absent and is supplied via the pre-computed `R9_pud_status.rds`.

7 Outcome Ascertainment: GERD and Oesophagitis

7.1 Two nested phenotypes

Both outcomes follow the project’s “ ≥ 2 diagnosis codes” specificity rule:

$$\begin{aligned}\text{has_gerd} &= \mathbb{1}\{\geq 2 \text{ GERD condition records}\}, \\ \text{has_esophagitis} &= \mathbb{1}\{\geq 2 \text{ GERD records named “esophagitis”}\}.\end{aligned}$$

Because every oesophagitis record is also a GERD record, `has_esophagitis` is a strict subset of `has_gerd`: any participant with two esophagitis-coded records necessarily has two GERD records. The two outcomes are nonetheless analyzed as *separate* binary outcomes on the full cohort (non-cases coded `FALSE`). Broad GERD is dominated by symptom-coded reflux and behaves partly as a disorder of gut–brain interaction; oesophagitis isolates the erosive subset with objective mucosal injury and is rarer. Requiring two records reduces misclassification from single “rule-out” codes and is the same heuristic used across the FGID portfolio.

7.2 Seed concepts and descendant expansion

A single BigQuery/Cohort Builder pull seeds two standard SNOMED concepts and expands to all standard, selectable descendants using the *All of Us* `cb_criteria` hierarchy pattern (the same mechanism used for the constipation outcome).

Concept ID	Standard concept name
318800	Gastroesophageal reflux disease
4223293	Gastroesophageal reflux disease with esophagitis

Table 6: GERD seed concepts. The `cb_criteria` join expands these to all standard descendants, so the pulled table carries many descendant `condition_concept_id` values beyond the two seeds.

The nested `SELECT` against `cb_criteria` converts the seed concepts to hierarchy `ids` (restricted to standard “rank-1” rows), then matches any condition concept whose materialized `path` lies on, under, or above those `ids`—capturing the entire descendant subtree. The outer query restricts to Fitbit-sharers (`has_fitbit = 1`) and joins concept names for readability. Concept IDs are seeds to be confirmed/expanded in the Cohort Builder before a production run; the recommended workflow searches “gastroesophageal reflux” and “reflux esophagitis,” selects the standard concept plus descendants, and pastes the full seed set into the `concept_id IN (...)` clause.

7.3 The descendant-filter subtlety

Because `R9_gerd_outcome.rds` is *already* a GERD-only pull (every row is a GERD descendant concept), `has_gerd` is computed over *all* rows of that table rather than by re-filtering on the two seed IDs. Re-filtering on only `c(318800, 4223293)` would silently drop every descendant-coded case and badly undercount GERD. The analysis therefore keeps a documented `GERD_CONCEPT_IDS` vector for transparency but leaves the in-analysis `filter()` commented out; a user who exports the *full* descendant list may paste it in and enable the explicit filter.

7.4 Identifying the oesophagitis subset

Within the pulled GERD rows, oesophagitis records are flagged by a concept-name match, `grepl("esophagitis|oesophagitis", standard_concept_name, ignore.case = TRUE)`. This robustly captures “GERD with esophagitis” and its descendants even if the exact descendant concept IDs vary across releases, and it is applied identically in the outcome-upload preview and in all three analysis files. A user who wishes to capture reflux/erosive esophagitis concepts that are *not* descendants of the two GERD seeds can add those seed IDs to the pull (documented in the upload script).

7.5 Case-definition variants

Variable	Rule	Purpose
<code>has_gerd</code>	≥ 2 GERD records	Primary broad outcome.
<code>has_esophagitis</code>	≥ 2 esophagitis-named records	Primary erosive outcome.
<code>has_gerd_dates</code>	≥ 2 records on <i>distinct</i> dates	Specificity sensitivity check.
<code>has_post_fitbit_gerd</code>	first GERD ≥ 180 days after first Fitbit, then ≥ 2 codes	Temporal ordering; reverse-causation guard.
<code>severe_gerd_binary</code>	GERD <i>and</i> ≥ 2 post-diagnosis PPI/ H_2 RA dates	Optional medication-treated (severe) GERD.

Table 7: Outcome and case-definition variants.

8 Exposure Definitions

8.1 Sleep metrics

For each retained night the pipeline keeps minutes asleep, in bed, awake, restless, and in the deep, light, and REM stages. Sleep efficiency is derived per night as

$$\text{efficiency} = \frac{\text{minutes asleep}}{\text{minutes in bed}}, \quad \text{defined only when minutes in bed} > 0.$$

Nightly values are averaged within participant to yield stable exposures (`avg_min_asleep`, `avg_min_in_bed`, `avg_min_awake`, `avg_min_restless`, `avg_min_deep`, `avg_min_light`, `avg_min_rem`, `avg_sleep_efficiency`) with `n_valid_nights`. Each averaged metric is discretized into cohort quartiles with `ntile(..., 4)` and stored as a four-level factor with **Q1 as reference**. Modeling quartiles rather than raw minutes sidesteps linearity assumptions, is robust to outliers, and yields interpretable dose-response patterns. The sleep-cleaning cascade retains main sleep only, keeps nights with minutes asleep in $(0, 1440]$, excludes participants with $\geq 30\%$ of nights below 240 minutes (4 hours), and requires ≥ 180 valid nights.

8.2 Activity metrics

A *valid activity day* has ≥ 10 wear hours and between 100 and 45,000 steps (pre-filtered in the intraday-steps table); participants need ≥ 30 valid days. On those days the pipeline computes

average daily steps (`avg_daily_steps`), activity-zone minutes (`avg_lightly_active_min`, `avg_fairly_active_min`, `avg_very_active_min`, `avg_total_active_min`, `avg_sedentary_min`), average daily wear hours, and, when available, the corrected maximum daily heart rate (`avg_daily_max_hr_minute_all_days`). Each is quantitized (Q1 reference); the max-HR exposure is guarded and omitted if its pre-built summary is absent.

Exposure set	Metric	Quartile variable
Sleep (8)	minutes asleep / in bed / awake / restless; deep; light; REM; efficiency	<code>min_asleep_quartile</code> , <code>min_in_bed_quartile</code> , <code>min_awake_quartile</code> , <code>min_restless_quartile</code> , <code>min_deep_quartile</code> , <code>min_light_quartile</code> , <code>min_rem_quartile</code> , <code>sleep_efficiency_quartile</code>
Activity (7)	daily steps; lightly/fairly/very/total active minutes; sedentary minutes; corrected max HR	<code>step_quartile</code> , <code>lightly_active_quartile</code> , <code>fairly_active_quartile</code> , <code>very_active_quartile</code> , <code>total_active_quartile</code> , <code>sedentary_quartile</code> , <code>max_hr_minute_all_days_quartile</code>

Table 8: The fifteen quartile exposures. The combined analysis models all fifteen (each separately and in mutually-adjusted pairs).

8.3 Sleep×activity overlap sensitivity set

A stricter sensitivity population retains only sleep-nights that coincide with a valid activity day (a same-day intraday-steps record), yielding `sleep_activity_overlap_df` and a parallel summary. This isolates participants with concurrently well-measured sleep and activity.

9 Combined Exposure Design

9.1 Intersection cohort

The combined analysis is defined on participants who satisfy *both* exposure-eligibility thresholds: ≥ 180 valid sleep nights *and* ≥ 30 valid activity days, plus EHR-sharing and age ≥ 18 . This is a genuine intersection (an `inner_join` of the two eligible sets), so its sample is smaller than either single-exposure cohort but every participant contributes both a complete sleep profile and a complete activity profile.

9.2 Covariate-timing anchor

For the single-exposure analyses, covariate windows are anchored to the first Fitbit date of the relevant stream (first sleep date for the sleep file; first activity date for the activity file). For the combined analysis, a participant may have started the two streams on different days, so covariate timing is anchored to the *earliest* Fitbit date across streams,

```
first_fitbit_date = min(first_sleep_date, first_activity_date),
```

computed element-wise with `pmin(..., na.rm = TRUE)`. All pre-Fitbit covariate windows (comorbidities, medications, BMI) are then defined consistently regardless of which stream began first.

9.3 Two model families

For each outcome the combined file fits: **(A)** one adjusted model per exposure, for all fifteen exposures, via the shared engine; and **(B)** *mutually-adjusted* models that include one sleep quartile and one activity quartile simultaneously, so each exposure is adjusted for the other domain in addition to the full covariate set. For a sleep exposure S , an activity exposure A , and covariates $\mathbf{Z}$,

$$\text{logit}(\Pr[Y = 1]) = \beta_0 + \sum_{j=2}^4 \beta_j^S \mathbb{I}[S=Q_j] + \sum_{j=2}^4 \beta_j^A \mathbb{I}[A=Q_j] + \sum_k \gamma_k Z_k.$$

Four primary sleep×activity pairs are examined: minutes-asleep×steps, efficiency×very-active, deep-sleep×total-active, and REM×sedentary. These pairings probe whether a sleep association survives adjustment for activity and vice versa.

10 Cohort Construction and Eligibility

10.1 Eligibility gates

Three participant-level gates apply to every configuration: age ≥ 18 at the anchoring first-Fitbit date; shared EHR data (present in the condition *or* measurement table); and shared Fitbit data, implied by meeting the stream-specific validity threshold (≥ 180 sleep nights, ≥ 30 activity days, or both for the combined cohort). Sleep cleaning additionally applies the main-sleep, valid-minute, and short-sleep rules of the previous section.

10.2 CONSORT accounting

Each analysis file reproduces a CONSORT-style waterfall counting unique participants surviving each step, from “any Fitbit record” through the final analytic cohort, recording the number excluded and the reason at each step.

Step	Sleep cohort criterion (kept)	Activity cohort criterion (kept)
0	Any Fitbit sleep record	Any Fitbit activity record
1	Main sleep & valid minutes (1–1440)	Valid day: ≥ 10 wear h, 100–45000 steps
2	$< 30\%$ nights with $< 4\text{h}$ sleep	—
3	≥ 180 valid nights	≥ 30 valid days
4	Shared any EHR data	Shared any EHR data
5	Age ≥ 18 at first Fitbit	Age ≥ 18 at first Fitbit

Table 9: CONSORT steps for the sleep and activity cohorts. The combined cohort is the intersection of the two final rows.

11 Rationale for Eligibility and Cleaning Thresholds

Every numeric threshold in the pipeline is a deliberate trade-off between measurement quality and sample size, chosen to match the antecedent IBS and constipation studies so that results are comparable across the portfolio. Table 10 collects them.

The sleep and activity coverage thresholds differ (≥ 180 nights vs ≥ 30 days) because they inherit from two separate parent studies and because sleep and step streams have different missingness profiles; this asymmetry is why the combined cohort is a genuine intersection rather than a simple union, and why its sample is the smallest of the three.

Rule	Threshold	Rationale
Main sleep only	<code>is_main_sleep</code>	Naps have different architecture and would bias nightly stage means.
Valid sleep minutes	(0, 1440]	Zero or > 24h asleep is physiologically impossible (device error).
Short-sleep exclusion	< 30% nights < 4h	Removes device-wear artifacts and extreme populations without discarding occasional short nights.
Sleep coverage	≥ 180 nights	About six months of data stabilizes long-run nightly means; shorter windows are noisier.
Valid activity day	≥ 10 wear h; 100–45000 steps	Ensures the device was worn most of the day and step counts are plausible (excludes non-wear and sensor spikes).
Activity coverage	≥ 30 valid days	Balances a stable activity average against retaining enough participants.
BMI window	± 180 days	Keeps adiposity contemporaneous with the exposure window.
Comorbidity rule	≥ 2 distinct pre-Fitbit dates	Enforces temporality and specificity (single stray codes excluded).
Medication rule	≥ 2 dates in prior 12 months	Captures sustained rather than incidental use.
Age gate	≥ 18 at first Fitbit	Adult population; consistent denominator.

Table 10: Eligibility and cleaning thresholds and their justifications.

12 Covariates

12.1 Temporal covariate rule

All EHR-derived comorbidities are ascertained *prior to* the anchoring first-Fitbit date and require ≥ 2 diagnoses on *distinct* dates, implemented by one helper, `create_covariate_status()`. Requiring two distinct dates before the exposure window enforces both temporality (the confounder precedes the exposure) and specificity (single stray codes do not qualify). This matches the functional-dyspepsia and constipation covariate builds, preserving cross-study consistency.

12.2 Covariate inventory

12.3 Medication, BMI, and GI covariate details

Medication “use” requires ≥ 2 distinct exposure dates within the 12 months preceding the first Fitbit date, capturing sustained rather than incidental use. BMI uses a within-person median (robust central tendency) and, in the single-exposure files, the single measurement closest to the first Fitbit date, preferring prior measurements and nulling values more than 180 days away. Two gastrointestinal covariates are added because they are clinically entangled with GERD: peptic ulcer disease (`has_pud`, shared acid pathophysiology and PPI treatment) from `R9_pud_status.rds`, and IBS (`has_ibs`), rebuilt *as a covariate* from `R9_condition_df` using the same pre-Fitbit, ≥ 2 distinct-dates rule so that GERD and IBS can co-occur while the temporal rule stays consistent.

13 The Shared Modeling Engine and Descriptive Statistics

13.1 Engine functions

Every configuration is executed by `GERD Analysis Helpers R9.R`. Its exported objects and functions are:

Domain	Variables	Source / definition
Demographics	age (+category), sex at birth, race, ethnicity	R9_person_df; collapsed.
Socio-economic	education, income	R9_survey_df items, collapsed.
Behavioral	smoking (ever ≥ 100 cigarettes), alcohol (Likert 0–5/day)	R9_survey_df.
Anthropometric	median BMI; nearest BMI within ± 180 days	R9_measurement_df (concept 3038553).
Psychiatric / cardiometabolic	depression, anxiety, diabetes, hypertension, heart failure, MI, stroke, COPD, sleep apnea	R9_condition_df via create_covariate_status().
Comorbidity burden	Charlson Comorbidity Index (continuous & categorical)	cci_scored_sleep.rds / R9_cci_scored.rds.
Medications	beta-blockers, calcium-channel blockers, stimulants, antidepressants, antipsychotics, anxiolytics, hypnotics	R9_drug_df; ≥ 2 dates within 12 months before first Fitbit.
GI (GERD-specific)	peptic ulcer disease (has_pud); IBS (has_ibs)	R9_pud_status.rds; IBS built from R9_condition_df as a covariate.

Table 11: Covariate inventory. The final row is the GERD-specific addition relative to the IBS study.

Condition	Concept ID	Condition	Concept ID
Depression	440383	Stroke	381316
Anxiety	441542	COPD	255573
Diabetes	201820	Sleep apnea (OSA)	313459
Hypertension	316866	IBS (covariate)	75576, 4234788, 4261072, 4057826
Heart failure	316139	Peptic ulcer disease	4318534, 4134146, 4265600, 4027663
Myocardial infarction	4329847		

Table 12: Condition concept IDs used to build covariates.

- GERD_OUTCOMES — the two outcomes and their factor labels; SLEEP_EXPOSURES and ACTIVITY_EXPOSURES; GERD_ADJ_COVARS (adds has_pud, has_ibs); and GERD_LABELS (master labels).
- prep_modeling_df(df, exposures) — builds collapsed factor covariates (race, ethnicity, sex, education, income, alcohol, smoking, Charlson category) and relevels each quartile exposure to Q1→Q4 with Q1 reference.
- descriptive_table(), univariate_table(), run_models_for_outcome(), stack_results(), run_diagnostics().
- analyze_all_outcomes(modeling_df, exposures) — the driver that loops both outcomes, producing descriptive, univariable, multivariable, and diagnostic output for each.

Each analysis file builds its one-row-per-participant frame, creates integer quartiles, calls prep_modeling_df(), then analyze_all_outcomes(). Centralization guarantees that the sleep, activity, and combined analyses — and both outcomes within each — are estimated by identical code.

13.2 Descriptive statistics

For each outcome a gtsummary::tbl_summary table is stratified by the outcome (labeled “No GERD”/“GERD” or “No oesophagitis”/“Oesophagitis”), reporting the exposure quartiles, continuous exposure metrics, demographics, behavioral and socio-economic factors, the Charlson index, all comorbidities, PUD, IBS, medications, and median BMI. Chi-square tests are attached for the collapsed categorical demographics, Benjamini–Hochberg FDR q -values are added, and

emphasis is placed on q rather than raw p . Race, ethnicity, sex, education, and income are collapsed to protect small cells and improve legibility.

14 Univariable and Multivariable Analysis

14.1 Univariable models

Each candidate predictor is regressed on the outcome individually via `gtsummary::tbl_uvregression` with `family = binomial` and exponentiated coefficients, yielding unadjusted odds ratios with 95% CIs for every exposure and covariate, including the new `has_pud` and `has_ibs` terms.

14.2 Multivariable models

Because the exposures within a set are collinear, a *separate* adjusted model is fit for each. For exposure X (a quartile factor with reference Q1) and covariate vector $\mathbf{Z}$,

$$\text{logit}(\Pr[Y = 1]) = \beta_0 + \beta_2\mathbb{I}[X=Q_2] + \beta_3\mathbb{I}[X=Q_3] + \beta_4\mathbb{I}[X=Q_4] + \sum_k \gamma_k Z_k,$$

so e^{β_j} is the adjusted odds ratio for quartile j versus Q1. The adjustment set is

```
age_cat, sex, race, ethnicity, education, income, alcohol, smoking, median_bmi,
has_depression, has_anxiety, has_pud, has_ibs, cci_cat.
```

For each exposure the engine fits a *full* model and a *backward* model selected by AIC (`MASS::stepAIC`) in which the exposure is forced to remain (the selection lower scope is `~ exposure`). Each model reports the complete-case N , the ROC AUC via `pROC`, and variance inflation factors via `car`; results are rendered as merged full-vs-backward tables and stacked across exposures. All six configurations reuse this exact procedure.

14.3 Combined mutually-adjusted models

The combined file additionally fits the mutually-adjusted sleep+activity models of Section 6.3, so that, for each outcome, the sleep and activity contributions are estimated in a single equation and each is adjusted for the other domain plus the full covariate set.

14.4 Assumption diagnostics

A dedicated diagnostics routine interrogates: **(1)** independence (no duplicate `person_id`); **(2)** multicollinearity (generalized VIF > 2 or VIF > 5); **(3)** separation/sparsity (logit standard errors > 5); **(4)** influential points (Cook's distance against $4/N$); and **(5)** linearity in the logit for continuous covariates (Box–Tidwell $x \log x$ term for median BMI). It runs for both outcomes of every configuration.

15 The Six Analysis Configurations in Detail

The design crosses two outcomes with three exposure sets. The six resulting configurations are summarized in Table 13 and described individually below. Configurations sharing a column are produced by the same analysis file, because `analyze_all_outcomes()` loops both outcomes over a single modeling data frame.

ID	Outcome	Exposure set	Cohort	File
C1	Broad GERD	Sleep (8 quartiles)	≥ 180 nights	Sleep and GERD
C2	Oesophagitis	Sleep (8 quartiles)	≥ 180 nights	Sleep and GERD
C3	Broad GERD	Activity (7 quartiles)	≥ 30 days	Activity and GERD
C4	Oesophagitis	Activity (7 quartiles)	≥ 30 days	Activity and GERD
C5	Broad GERD	Combined (15 + pairs)	Intersection	Activity Sleep and GERD
C6	Oesophagitis	Combined (15 + pairs)	Intersection	Activity Sleep and GERD

Table 13: The six analysis configurations in detail.

15.1 C1 — Sleep and broad GERD

The reference configuration. On the ≥ 180 -night sleep cohort, each of the eight sleep quartiles is regressed on `has_gerd` unadjusted and then adjusted for the full covariate set, with a full-vs-backward pairing per exposure. This is the direct GERD analogue of the project’s IBS sleep study and the primary test of the “better sleep, lower reflux odds” hypothesis.

15.2 C2 — Sleep and oesophagitis

Identical machinery, outcome `has_esophagitis`. Because the erosive subset is rarer, case counts per quartile are smaller and confidence intervals wider; the separation diagnostic is especially relevant. Directional agreement with C1 argues that any sleep association acts on mucosal injury and not only on symptom reporting.

15.3 C3 — Activity and broad GERD

On the ≥ 30 -valid-day activity cohort, the seven activity quartiles (steps, four activity-intensity bands, sedentary minutes, and corrected max heart rate) are modeled against `has_gerd`. Sedentary time and low step counts are hypothesized to associate with higher odds; the corrected max-HR exposure provides a cardiovascular-fitness proxy.

15.4 C4 — Activity and oesophagitis

The activity exposures against the erosive subset, again precision-limited by case count. Comparison with C3 mirrors the C1/C2 contrast in the activity domain.

15.5 C5 — Combined exposures and broad GERD

On the intersection cohort, all fifteen exposures are modeled separately (as in C1 and C3 but on the common sample), and the four mutually-adjusted sleep $\times$ activity pairs are fit so that each domain is adjusted for the other. This isolates independent contributions: a sleep association that survives adjustment for steps, and an activity association that survives adjustment for sleep duration, are more credible than either single-domain estimate.

15.6 C6 — Combined exposures and oesophagitis

The most demanding configuration: the erosive outcome on the (smallest) intersection cohort with mutually-adjusted models. It is interpreted cautiously and primarily for consistency of direction with C5, C2, and C4.

16 The Logistic Model in Depth

16.1 Specification and estimation

Each adjusted model is a binomial generalized linear model with the canonical logit link. Writing $p_i = \Pr[Y_i = 1 \mid X_i, Z_i]$ for participant i , the model is

$$\log \frac{p_i}{1 - p_i} = \beta_0 + \sum_{j=2}^4 \beta_j \mathbb{I}[X_i = Q_j] + \sum_k \gamma_k Z_{ik},$$

estimated by maximum likelihood via iteratively reweighted least squares (R's `glm`). The quartile exposure enters as a three-parameter contrast against Q1, so the model makes no assumption of linearity across quartiles; a U -shaped or threshold relationship is representable.

16.2 Interpretation of the quartile contrasts

The exponentiated coefficient e^{β_j} is the odds ratio comparing quartile j to Q1, adjusted for all covariates. Reporting all three contrasts (Q_2, Q_3, Q_4 vs Q_1) reveals the *shape* of the exposure–outcome relationship rather than compressing it to a single slope. A monotone trend in $e^{\beta_2}, e^{\beta_3}, e^{\beta_4}$ suggests dose–response; a non-monotone pattern (for example, an elevated top and bottom quartile) suggests a non-linear or mixed-population effect worth further study.

16.3 Variable selection and the backward model

Alongside each full model, a backward-elimination model is selected by the Akaike Information Criterion (`MASS::stepAIC, direction = "backward"`) with the exposure *forced* to remain (the lower scope is `~ exposure`). The backward model is a parsimony check: if the exposure's estimate is stable between the full and backward models, the association is not an artifact of a particular covariate set. AIC balances fit against the number of parameters, discouraging overfitting in the smaller (oesophagitis, combined) cohorts.

16.4 Discrimination and multicollinearity

Each model reports the area under the receiver-operating-characteristic curve (AUC) via `pROC` as a discrimination summary, and generalized variance inflation factors (GVIF) via `car`. For factor covariates the reported quantity is $\text{GVIF}^{1/(2\,df)}$, comparable across terms of differing degrees of freedom; values above 2 (or ordinary VIF above 5) flag problematic collinearity. Because the exposures within a set are correlated, they are deliberately *not* entered together in the single-exposure models; the combined mutually-adjusted models pair only one sleep and one activity exposure at a time to keep collinearity manageable.

17 Model Assumptions and Remediation

The diagnostics routine is not merely descriptive; each check is tied to a concrete remediation should it fail. Table 14 lists the assumption, the test, what a violation implies, and the recommended response.

The exposure quartiles are categorical and therefore exempt from the linearity-in-the-logit requirement; only genuinely continuous covariates (median BMI, and Charlson score if used continuously) are tested. The separation check is most consequential for the rarer oesophagitis and treated-GERD outcomes, where a covariate level may perfectly predict non-case status; the routine surfaces this as an abnormally large standard error rather than letting it pass silently.

Assumption	Test	Violation implies	Remediation
Independence	Duplicate <code>person_id</code> check	Correlated observations, understated SEs	Deduplicate joins; one row per participant.
No severe multi-collinearity	$\text{GVIF}^{1/(2df)} > 2$ / $\text{VIF} > 5$	Unstable, inflated coefficients	Drop or combine collinear covariates; rely on the backward model.
No perfect separation	Logit SE > 5	Non-identifiable coefficients	Collapse sparse categories; consider penalized (Firth) logistic regression.
No overly influential points	Cook's distance vs $4/N$	A few rows drive the fit	Inspect flagged rows; sensitivity refit excluding them.
Linearity in the logit	Box–Tidwell $x \log x$ for BMI	Continuous covariate mis-specified	Categorize BMI or add a spline/quadratic term.

Table 14: Logistic-regression assumptions, diagnostics, and remediation.

18 Data Engineering, Quality Control, and Statistical Considerations

18.1 One-row-per-participant guarantee

Every derived covariate table is reduced to one row per `person_id` before joining, and the diagnostics assert no duplicate identifiers in the modeling frame. Eligibility is defined once (an `inner_join` to the filtered exposure summary), and covariates are attached by `left_join` so they never expand the row count.

18.2 Missing data

Binary EHR indicators (conditions, medications, PUD, IBS, and the outcomes) are coalesced from NA to FALSE: in EHR phenotyping the absence of a qualifying record means the participant does not meet the phenotype. Charlson score is coalesced to 0 when no components precede the exposure window. Survey/measurement covariates (education, income, alcohol, BMI) are *not* imputed; an explicit “Unknown/Missing” level is carried through and reported, so differential missingness by outcome is visible. Each adjusted model uses complete cases over its own variable set, so the effective N is reported per model and can differ across exposures.

18.3 Multiple comparisons

With fifteen exposures (each with several quartile contrasts) against two outcomes across three exposure sets, multiplicity is substantial. Descriptive and treated-GERD tables apply Benjamini–Hochberg FDR q -values via `add_q(method = "fdr")`, and the primary multivariable emphasis is on effect sizes (aOR) and confidence intervals under a fixed, pre-specified adjustment set rather than on dichotomous significance.

18.4 Confounding structure

Each covariate blocks a plausible backdoor path: adiposity (BMI) raises intra-abdominal pressure and also shortens/fragments sleep and reduces activity; age, sex, race, ethnicity, and socio-economic status pattern both behavior and health-care coding; smoking and alcohol relax the lower esophageal sphincter and degrade sleep; depression and anxiety alter sleep and amplify symptom reporting; IBS and PUD co-travel with GERD; the Charlson index captures residual multimorbidity; and sleep-relevant medications affect measured sleep. Mediators of the

exposure→GERD path (acid-suppression therapy) are excluded from the primary adjustment set to avoid over-adjustment and used only to define the treated-GERD sensitivity outcome.

18.5 Sample size and precision

No fixed sample size is imposed; precision is governed by the number of *cases*. Oesophagitis, being the erosive subset, is rarer than broad GERD, and the treated-GERD sensitivity (a further subset) is the most precision-limited. The combined intersection cohort is smaller than either single-exposure cohort. Sparse covariate×outcome cells are exactly what the separation diagnostic is designed to catch, and backward-AIC can drop unstable terms while the exposure is retained.

19 Sensitivity Analyses

1. **Distinct-date GERD** (`has_gerd_dates`) — the stricter “ ≥ 2 distinct dates” phenotype, testing robustness to same-day code inflation.
2. **Post-Fitbit temporal GERD** (`has_post_fitbit_gerd`) — restricts cases to those whose first GERD code appears ≥ 180 days *after* the first Fitbit date, so the measured exposure precedes the outcome’s EHR appearance. This is the strongest available guard against reverse causation, which is a genuine concern because nocturnal reflux fragments sleep.
3. **Sleep×activity overlap** — repeats analyses on the concurrently well-measured subset.
4. **Medication-treated (“severe”) GERD** (`severe_gerd_binary`) — an optional outcome flagging participants who both meet the GERD phenotype *and* have ≥ 2 PPI/ H_2 RA prescription dates beginning at least one day after their first GERD diagnosis, an EHR proxy for clinically managed disease. It mirrors the IBS study’s “severe IBS” construction but substitutes acid-suppression pharmacotherapy.

19.1 The pharmacologics pull

The treated-GERD analysis depends on `GERD Pharmacologics Data Upload.R`, which retrieves drug exposures whose ancestor concept is a PPI (21600095) or an H_2 RA (21600096) using the `concept_ancestor` double-join (the same structure as the IBS drug pull), restricted to Fitbit-sharers. The whole treated-GERD section is guarded: it runs only if `R9_gerd_drug_df.rds` exists, so the primary pipeline executes cleanly whether or not the optional pull has been performed.

20 Relationship to the IBS and Constipation Templates

The suite mirrors the project’s IBS sleep study, constipation activity study, and constipation outcome upload, but global find-and-replace was insufficient; whole blocks were removed, added, or corrected.

Every GERD-specific artifact is `R9_`-prefixed or `_gerd`-suffixed, so running the GERD pipeline never overwrites an IBS or constipation artifact even though the underlying sleep-cleaning and activity-summary steps are numerically identical.

21 Reproducibility and Execution

21.1 Execution order

1. In the workspace, verify the GERD/oesophagitis seed concepts in the Cohort Builder, then run `GERD Outcome Data Upload R9.R → R9_gerd_outcome.rds`.

Action	Detail
Swapped	<code>has_ibs</code> → <code>has_gerd</code> and added <code>has_esophagitis</code> ; outcome labels; output filenames → <code>R9_final_analysis*_gerd_df.rds</code> .
Swapped source	IBS rows from <code>dataset_32338507_condition_df</code> → the dedicated <code>R9_gerd_outcome</code> pull.
Removed	All IBS-subtype logic (subtype flags, most-frequent/tie-break assignment, mutually-exclusive <code>has_ibs_c/d/m/general</code> , multinomial models); the standalone severe-IBS block was replaced by the optional, guarded treated-GERD section.
Added	<code>has_esophagitis</code> as a second outcome; <code>has_pud</code> and <code>has_ibs</code> covariates; the activity and combined analyses; the mutually-adjusted sleep+activity models; and the shared modeling engine.
Fixed	Consistent <code>R9_*</code> inputs throughout; one canonical <code>first_fitbit_sleep_date_df</code> / <code>first_fitbit_date_df</code> and <code>sleep_clean_df</code> , resolving the naming mismatches present in the IBS script.

Table 15: IBS/constipation → GERD/oesophagitis transformation.

2. Ensure GERD Analysis Helpers R9.R is in the working directory (each analysis file `source()`s it).
3. Run any of Sleep and GERD Analysis File R9.R, Activity and GERD Analysis File R9.R, and Activity Sleep and GERD Analysis File R9.R — each produces both outcomes.
4. (Optional) Run GERD Pharmacologics Data Upload.R → `R9_gerd_drug_df.rds`, then re-run the sleep analysis to activate the treated-GERD section.

21.2 Software environment

The pipeline uses `tidyverse`, `gtsummary`, `broom`, `lubridate`, `MASS`, `car`, `pROC`, `forcats`, and `rlang`, plus `bigquery` for the pulls. The helper installs missing packages on first load. Because all modeling lives in the helper, a change there propagates to all six configurations at once.

22 Anticipated Results and Interpretation Framework

Consistent with the FGID portfolio’s hypothesis, we anticipate that higher sleep quantity and quality and higher physical activity will be associated with *lower* adjusted odds of both GERD and oesophagitis, and that markers of fragmented sleep or high sedentary time will be associated with *higher* odds. An aOR of, say, 0.85 for Q4 versus Q1 of `min_asleep_quartile` means that, holding the adjustment set fixed, participants in the highest sleep-duration quartile have 15% lower odds of the outcome than those in the lowest. A monotone decline across Q2→Q3→Q4 constitutes a dose–response pattern.

Three internal-consistency checks structure interpretation. First, *concordance across outcomes*: if an exposure association is directionally consistent for broad GERD and for the erosive oesophagitis subset, it is less likely to be an artifact of symptom-reporting behavior and more likely to reflect a mucosal-injury pathway. Second, *concordance across exposure domains*: agreement between the sleep-only, activity-only, and mutually-adjusted combined estimates indicates that neither domain is merely a proxy for the other. Third, *temporal concordance*: agreement between the cross-sectional primary estimate and the post-Fitbit temporal definition is the key guard against reverse causation. Because the design is cross-sectional, the primary estimates remain associations, not causal effects.

Exposure (high vs low)	Predicted aOR	Primary mechanism
More minutes asleep	< 1 (protective)	More clearance, less arousal, lower sensitivity.
Higher sleep efficiency	< 1	Consolidated sleep, less supine low-clearance time.
More deep / REM sleep	< 1	Restorative sleep, autonomic balance.
More minutes awake / restless	> 1 (harmful)	Fragmentation, sympathetic arousal.
More daily steps / active minutes	< 1	Weight control, vagal tone, motility.
More sedentary minutes	> 1	Adiposity, postprandial recumbency.
Higher corrected max HR	< 1 (fitness proxy)	Autonomic/cardiovascular fitness.

Table 16: Pre-specified directional predictions for the exposure–outcome associations (both outcomes). Predictions are expected to be directionally similar for GERD and oesophagitis, with acid-contact pathways potentially stronger for the erosive subset and perceptual pathways stronger for broad GERD.

23 Threats to Validity

Beyond the item-level limitations below, it is useful to organize the principal threats to validity by epidemiologic type and to state how the design addresses each.

23.1 Selection bias

The cohort is restricted to participants who own a Fitbit, share it with *All of Us*, wear it enough to meet the coverage thresholds, and share EHR data. Each filter can associate with both exposure and outcome (for example, health-motivated device use). The design cannot remove this but makes it transparent through the CONSORT waterfall, and adjusts for socio-economic status, which is a principal driver of device access.

23.2 Information (measurement) bias

The exposure is objective (device-recorded) but the outcome is EHR-coded and therefore subject to detection and coding variation: participants with more health-care contact are both more likely to be coded GERD and may differ in sleep and activity. Conditioning on the Charlson index and comorbidities partially blocks this “surveillance” pathway. The ≥ 2 -record and ≥ 2 -distinct-date rules reduce false-positive coding at the cost of sensitivity.

23.3 Confounding

Addressed by the pre-specified adjustment set (Section 9.2), the pre-Fitbit temporal covariate rule (so covariates precede exposure), and the mutually-adjusted combined models (so each exposure domain is adjusted for the other). Residual confounding from unmeasured factors (diet, meal timing, caffeine, hiatal hernia) remains possible and is stated as a limitation.

23.4 Reverse causation

The most specific threat for reflux: nocturnal GERD fragments sleep, so a cross-sectional sleep–GERD association could run outcome→exposure. The post-Fitbit temporal outcome (`has_post_fitbit_gerd`), which requires the first GERD code to appear well after the exposure window, is the primary guard; concordance between it and the primary estimate is the key internal check.

23.5 External validity

Findings generalize to adult, EHR-sharing, Fitbit-wearing *All of Us* participants with sufficient wear time — a group that is younger, more affluent, and more health-engaged than the general population. Estimates should not be extrapolated to non-wearers without caution.

24 Limitations

- **Cross-sectional design.** Temporality is only partially addressed; primary estimates cannot establish causation, and reflux itself disrupts sleep.
- **EHR phenotyping.** Diagnosis-code outcomes miss undiagnosed or over-the-counter-managed reflux and vary with coding practice; the ≥ 2 -record rule trades sensitivity for specificity.
- **Name-based oesophagitis identification.** Defining the erosive subset by concept-name match captures “GERD with esophagitis” but may miss erosive-esophagitis concepts coded outside the GERD hierarchy unless their seeds are added in the Cohort Builder.
- **Wearable selection.** Fitbit users are not representative of the general population.
- **Residual confounding.** Diet, meal timing, caffeine, body position, and hiatal hernia remain unmeasured; hiatal hernia would require a new condition pull.
- **Concept-set completeness.** GERD and drug-class concept IDs are seeds requiring Cohort Builder verification; the descendant-expansion design (Section 7.3) mitigates but does not eliminate this.
- **Covariate-timing heterogeneity.** `has_pud` is loaded from a file computed on the activity-based first-Fitbit date; the small resulting timing difference is a minor, documented caveat.

25 Future Directions and Extensions

The suite is deliberately extensible. Several extensions follow naturally from the shared-engine design, each requiring only a new modeling data frame or a new outcome column while reusing the same functions.

- **Dose–response modeling.** Replace the quartile factors with restricted cubic splines on the continuous metrics to estimate smooth exposure–response curves, using the Box–Tidwell result to decide where non-linearity matters.
- **Effect-measure modification.** Add interaction terms (for example, $\text{sleep} \times \text{BMI}$, or $\text{sleep} \times \text{sex}$) to test whether the reflux association differs across subgroups.
- **Erosive-severity gradient.** Extend the oesophagitis phenotype toward a graded outcome (reflux esophagitis, stricture, Barrett metaplasia) to test a monotone exposure–severity relationship.
- **Fatigue and symptom stratification.** Incorporate the fatigue-questionnaire data used in the sibling IBS work to stratify or adjust the sleep analyses.
- **Longitudinal follow-up.** Move from cross-sectional prevalence to incident GERD after the Fitbit window, strengthening the temporal argument beyond the current post-Fitbit sensitivity analysis.
- **Medication-response phenotyping.** Extend the treated-GERD flag to distinguish PPI-responsive from PPI-refractory disease, a clinically important axis of reflux heterogeneity.

- **Additional physiological exposures.** Heart-rate variability and resting heart rate, if available, would sharpen the autonomic-pathway hypothesis.

Because every extension plugs into `analyze_all_outcomes()` through the same helper, adding one does not perturb the six primary configurations, preserving the internal consistency that the architecture is designed to guarantee.

26 Integration Across the FGID Portfolio

This GERD/oesophagitis suite is the fourth member of a coordinated program that already covers irritable bowel syndrome, functional dyspepsia, and constipation, each relating the same Fitbit-derived sleep and activity exposures to a disorder of gut–brain interaction. Because the studies share eligibility rules, exposure quartilization, covariate definitions, and (now) a common modeling engine, their estimates are directly comparable. Table 17 places GERD within the portfolio.

Study	Outcome(s)	Distinctive features
IBS	<code>has_ibs</code> (+ subtypes)	Multinomial subtype models; severe-IBS medication subgroup.
Functional dyspepsia	<code>has_fd</code> (strict/broad)	Rome-based phenotypes; early-satiety extension.
Constipation	<code>has_constipation</code>	Simple binary outcome; heart-rate-correction activity analysis.
GERD / oesophagitis	<code>has_gerd;</code> <code>has_oesophagitis</code>	Two nested acid-related phenotypes; PUD and IBS covariates; sleep, activity, and combined exposures under one shared engine.

Table 17: The GERD/oesophagitis analysis within the functional-GI-disorder portfolio.

Cross-disorder comparison is scientifically informative. If, for example, poor sleep is associated with higher odds across all four disorders, that pattern supports a shared gut–brain mechanism rather than a reflux-specific one; if the association is strongest for the erosive oesophagitis phenotype, it points to an acid-contact pathway specific to the upper GI tract. The portfolio design — identical exposures, harmonized covariates, and a common engine — is what makes such comparisons valid rather than confounded by differing analytic choices. The GERD suite’s two nested outcomes add a within-disorder axis (symptom-coded vs erosive) that the other studies do not have, giving the portfolio its first built-in objective-injury comparator.

27 Conclusion

This report specifies a complete, reproducible suite relating objectively measured sleep and physical activity to two nested EHR-defined outcomes — broad GERD and erosive oesophagitis — in *All of Us* CDR v9. Six configurations (two outcomes × three exposure sets) are executed by a single shared modeling engine that guarantees statistical identity across the sleep, activity, and combined analyses. The suite faithfully mirrors the project’s IBS and constipation studies while correcting their naming inconsistencies, removing IBS-only subtype machinery, adding GERD-appropriate gastrointestinal covariates and a second erosive outcome, and extending the design to physical activity and to mutually-adjusted combined models. Together with distinct-date, temporal, overlap, and treated-disease sensitivity analyses, the design supports a robust, portfolio-consistent assessment of whether sleep and activity are associated with reflux disease and its erosive form.

A Appendix A: Output Inventory

Output .rds / object	Description
R9_gerd_outcome.rds	Raw GERD condition records (+ is_esophagitis flag).
R9_gerd_status.rds	has_gerd (≥ 2 records) per participant.
R9_esophagitis_status.rds	has_esophagitis (≥ 2 esophagitis-named records).
R9_gerd_status_dates.rds	has_gerd_dates (≥ 2 distinct dates).
R9_gerd_post_fitbit_sleep.rds	has_post_fitbit_gerd temporal flag (sleep anchor).
R9_final_analysis_sleep_gerd_df.rds	Sleep modeling frame (configs C1, C2).
R9_final_analysis_activity_gerd_df.rds	Activity modeling frame (configs C3, C4).
R9_final_analysis_activity_sleep_gerd_df.rds	Combined modeling frame (configs C5, C6).
R9_comorbidity_status_sleep_gerd_df.rds	Combined comorbidity indicators.
R9_ibs_covariate_status_gerd.rds	IBS covariate indicator.
R9_cci_covariates_sleep_gerd_df.rds	Charlson score and category.
R9_bmi_covariates_sleep_gerd_df.rds	BMI covariates.
R9_medication_flags_sleep_gerd.rds	Medication-use flags.
R9_rx_post_gerd_df.rds	Treated-GERD prescription flag (optional).

B Appendix B: Concept Reference

Purpose	Concept ID(s)	Note
GERD (outcome seeds)	318800; 4223293	Descendant-expanded.
Oesophagitis subset	4223293 (+ name match)	Erosive subset.
BMI measurement	3038553	Body mass index.
IBS (covariate)	75576; 4234788; 4261072; 4057826	General + subtypes.
Peptic ulcer disease	4318534; 4134146; 4265600; 4027663	Via R9_pud_status.
Proton pump inhibitors	21600095	Drug-class ancestor (verify).
H ₂ -receptor antagonists	21600096	Drug-class ancestor (verify).
Beta-blockers	21601664	Medication class.
Calcium-channel blockers	21601765	Medication class.
Stimulants	21604753	Medication class.
Antidepressants	21604686	Medication class.
Antipsychotics	21604490	Medication class.
Anxiolytics	21604600; 21604565	Medication class.
Hypnotics	21604635; 21604653; 21604685; 21604661	Medication class.

Table 19: Consolidated concept-ID reference.

C Appendix C: Selected Code Excerpts

```

1 # Broad GERD: ≥2 GERD records (computed over ALL rows of the GERD-only pull)
2 R9_gerd_status <- R9_gerd_outcome %>%
3   group_by(person_id) %>% summarise(has_gerd = n() ≥ 2, .groups = "drop")
4
5 # Oesophagitis (erosive) subset: ≥2 records whose concept name is esophagitis
6 R9_esophagitis_status <- R9_gerd_outcome %>%
7   filter(grepl("esophagitis|oesophagitis", standard_concept_name, ignore.case = TRUE)) %>%
8   group_by(person_id) %>% summarise(has_esophagitis = n() ≥ 2, .groups = "drop")

```

Listing 1: Two nested outcomes from one GERD pull (analysis files).

```

1 source("GERD Analysis Helpers R9.R")
2 modeling_df <- prep_modeling_df(final_analysis_df, SLEEP_EXPOSURES) # or ACTIVITY_EXPOSURES / union
3 # Runs descriptive + univariable + multivariable + diagnostics for BOTH
4 # has_gerd and has_esophagitis over every exposure:
5 results <- analyze_all_outcomes(modeling_df, exposures = SLEEP_EXPOSURES)

```

Listing 2: The shared driver: both outcomes, one exposure set.

```

1 run_models_for_outcome <- function(df, outcome_col, outcome_labels, exposures, covars, ...) {
2   d0 <- factor_outcome(df, outcome_col, outcome_labels)
3   fit_one <- function(exposure) {
4     d <- d0 %>% select(all_of(c(outcome_col, exposure, covars))) %>% drop_na()
5     m_full <- glm(as.formula(paste0(outcome_col, " ~ ",
6                                   paste(c(exposure, covars), collapse = " + "))), d, family = binomial())
7     m_step <- MASS::stepAIC(m_full, direction = "backward",
8                           scope = list(lower = as.formula(paste0("~ ", exposure)),
9                                       upper = formula(m_full)), trace = FALSE)
10    # AUC (pROC), VIF (car), gtsummary full-vs-backward merge ...
11  }
12  purrr::map(intersect(exposures, names(d0)), fit_one)
13 }

```

Listing 3: Per-exposure adjusted modeling inside the engine (abridged).

```

1 fit_combined_pair <- function(df, outcome_col, outcome_labels, sleep_expo, act_expo, covars) {
2   d <- df; d[[outcome_col]] <- factor(d[[outcome_col]], c(FALSE, TRUE), labels = outcome_labels)
3   d <- d %>% select(all_of(c(outcome_col, sleep_expo, act_expo, covars))) %>% drop_na()
4   glm(as.formula(paste0(outcome_col, " ~ ", sleep_expo, " + ", act_expo, " + ",
5                         paste(covars, collapse = " + "))), d, family = binomial()) # each domain adjusts for the other
6 }

```

Listing 4: Combined mutually-adjusted model: sleep + activity together.

D Appendix D: Glossary of Key Derived Variables

Variable	Meaning
first_fitbit_sleep_date	Earliest valid Fitbit sleep date; anchors sleep-file covariate windows.
first_fitbit_date	Earliest Fitbit date (activity file: activity stream; combined file: min of both streams).
age_at_fitbit_start / age_cat	Age and its category (<30, 30–44, 45–59, 60–74, 75+) at the first Fitbit date.
n_valid_nights / n_valid_days	Valid sleep nights / valid activity days (eligibility gates).
avg_min_asleep, ..., avg_sleep_efficiency	Participant-level mean of each nightly sleep metric.

Variable	Meaning
avg_daily_steps, avg_*_active_min, avg_sedentary_min	Participant-level mean daily steps and activity-zone minutes.
avg_daily_max_hr_minute_-all_days	Corrected mean of daily maximum heart rate.
min_asleep_quartile, ...	Cohort quartile (Q1 reference) of a sleep metric — a modeled exposure.
step_quartile, ..., max_hr_minute_all_days_quartile	Cohort quartile of an activity metric — a modeled exposure.
has_gerd	Primary broad outcome: ≥ 2 GERD records.
has_esophagitis	Primary erosive outcome: ≥ 2 esophagitis-named records.
has_gerd_dates	Sensitivity: ≥ 2 GERD records on distinct dates.
has_post_fitbit_gerd	Temporal sensitivity: first GERD ≥ 180 days after first Fitbit, then ≥ 2 codes.
severe_gerd_binary	Optional: GERD <i>and</i> ≥ 2 post-diagnosis PPI/ H_2 RA dates.
has_pud, has_ibs	GERD-specific gastrointestinal covariates.
has_depression ... has_sleep_apnea	Pre-Fitbit, ≥ 2 -distinct-date comorbidity indicators.
on_beta_blocker ... on_hypnotics	Sustained pre-exposure medication-use flags.
cci_score / cci_cat	Charlson Comorbidity Index, continuous and categorical (0, 1–2, 3–4, 5+).
median_bmi / bmi_closest	Within-person median BMI; nearest BMI within ± 180 days.
alcohol_likert_final	Daily-drinking Likert (0 = none through 5 = ≥ 10 /day).
smoking_binary	Ever smoked ≥ 100 cigarettes (1) vs not (0).
education_collapsed, income_collapsed	Collapsed socio-economic factors.
final_analysis_sleep_gerd_df	Sleep modeling frame (configs C1, C2).
final_analysis_activity_gerd_df	Activity modeling frame (configs C3, C4).
final_analysis_activity_sleep_gerd_df	Combined modeling frame (configs C5, C6).

E Appendix E: Results-Template Tables

The analyses are not yet run, so numeric cells are left as placeholders (“–”). These shells fix the intended reporting format; each analysis file produces exactly these tables (via `gtsummary::tbl_stack()` for both outcomes.

E.1 Adjusted odds ratios — sleep exposures (C1 broad GERD; C2 oesophagitis)

E.2 Adjusted odds ratios — activity exposures (C3 broad GERD; C4 oesophagitis)

E.3 Mutually-adjusted combined pairs (C5 broad GERD; C6 oesophagitis)

Sleep exposure (Q4 vs Q1)	C1: broad GERD		C2: oesophagitis	
	aOR	95% CI	aOR	95% CI
Minutes asleep	–	(–,–)	–	(–,–)
Minutes in bed	–	(–,–)	–	(–,–)
Minutes awake	–	(–,–)	–	(–,–)
Minutes restless	–	(–,–)	–	(–,–)
Deep sleep	–	(–,–)	–	(–,–)
Light sleep	–	(–,–)	–	(–,–)
REM sleep	–	(–,–)	–	(–,–)
Sleep efficiency	–	(–,–)	–	(–,–)

Table 21: Template: adjusted odds ratios (top vs bottom quartile) for the sleep exposures against both outcomes. Full tables also report Q2 and Q3 contrasts, backward models, N , and AUC.

Activity exposure (Q4 vs Q1)	C3: broad GERD		C4: oesophagitis	
	aOR	95% CI	aOR	95% CI
Daily steps	–	(–,–)	–	(–,–)
Lightly active min	–	(–,–)	–	(–,–)
Fairly active min	–	(–,–)	–	(–,–)
Very active min	–	(–,–)	–	(–,–)
Total active min	–	(–,–)	–	(–,–)
Sedentary min	–	(–,–)	–	(–,–)
Corrected max HR	–	(–,–)	–	(–,–)

Table 22: Template: adjusted odds ratios for the activity exposures against both outcomes.

E.4 Descriptive-table shell (per outcome)

F Appendix F: Per-File Walkthrough

This appendix summarizes, step by step, what each of the six deliverable scripts does.

F.1 GERD Outcome Data Upload R9.R

1. Pins the workspace CDR to v9 and constructs the `cb_criteria` descendant-expansion SQL for seeds 318800 and 4223293, restricted to Fitbit-sharers.
2. Exports the query to the workspace bucket, reads it back, and previews the captured GERD concepts.
3. Flags the oesophagitis subset by concept-name match and prints its size.
4. Saves `R9_gerd_outcome.rds` (with the `is_oesophagitis` flag).

F.2 GERD Analysis Helpers R9.R

Defines the outcome objects, exposure lists, adjustment covariates, and label lookup, and the functions `prep_modeling_df()`, `descriptive_table()`, `univariate_table()`, `run_models_for_outcome()`, `stack_results()`, `run_diagnostics()`, and `analyze_all_outcomes()`. It installs any missing packages on load and is sourced by all three analysis files.

F.3 Sleep and GERD Analysis File R9.R

1. Loads R9 inputs; computes the first sleep date and age; cleans sleep records; builds participant-level sleep summaries.
2. Builds both outcomes (`has_gerd`, `has_oesophagitis`) and the sensitivity definitions.

Sleep term	Activity term	C5: broad GERD		C6: oesophagitis	
		aOR _S	aOR _A	aOR _S	aOR _A
Minutes asleep (Q4)	Steps (Q4)	–	–	–	–
Efficiency (Q4)	Very active (Q4)	–	–	–	–
Deep sleep (Q4)	Total active (Q4)	–	–	–	–
REM sleep (Q4)	Sedentary (Q4)	–	–	–	–

Table 23: Template: mutually-adjusted models. aOR_S and aOR_A are the top-vs-bottom-quartile odds ratios for the sleep and activity terms, each adjusted for the other domain and the full covariate set.

Characteristic	Non-case	Case	<i>q</i> -value
N	–	–	
Age group, n (%)	–	–	–
Sex at birth, n (%)	–	–	–
Race / ethnicity, n (%)	–	–	–
Median BMI	–	–	–
Charlson index, n (%)	–	–	–
Depression / anxiety, n (%)	–	–	–
Peptic ulcer disease, n (%)	–	–	–
IBS, n (%)	–	–	–
Exposure quartiles, n (%)	–	–	–

Table 24: Template: descriptive table stratified by an outcome (produced for each of the six configurations), with Benjamini–Hochberg FDR *q*-values.

3. Builds behavioral, comorbidity, PUD, IBS, Charlson, BMI, medication, and socio-economic covariates on the sleep anchor.
4. Applies eligibility, assembles `R9_final_analysis_sleep_gerd_df`, creates integer sleep quartiles, sources the helper, and calls `analyze_all_outcomes()` for both outcomes.
5. Runs the optional treated-GERD block if the drug file exists.

F.4 Activity and GERD Analysis File R9.R

Parallel structure on the activity stream: first activity date and age; steps, activity-zone, wear, and corrected max-HR summaries; the same outcomes and covariates on the activity anchor; the ≥ 30 -valid-day eligibility; `R9_final_analysis_activity_gerd_df`; integer activity quartiles; and `analyze_all_outcomes()` for both outcomes.

F.5 Activity Sleep and GERD Analysis File R9.R

Builds both exposure streams, anchors covariates to the earliest Fitbit date, forms the intersection cohort, assembles `R9_final_analysis_activity_sleep_gerd_df` with both quartile sets, runs `analyze_all_outcomes()` over the union of exposures, and then fits the four mutually-adjusted sleep×activity pairs for both outcomes.

F.6 GERD Pharmacologics Data Upload.R

Pulls PPI (ancestor 21600095) and H_2RA (ancestor 21600096) exposures via the `concept_ancestor` double-join for Fitbit-sharers and saves `R9_gerd_drug_df.rds`, enabling the treated-GERD sensitivity analysis.

G Appendix G: Reporting and Reproducibility Checklist

Item	Where addressed
Study design	Retrospective cross-sectional; Sections on design and configurations.
Setting / data source	<i>All of Us</i> Controlled Tier CDR v9; Data Sources.
Participants / eligibility	CONSORT and threshold-rationale sections.
Exposures	Sleep and activity quartile definitions; Exposure Definitions.
Outcomes	GERD and oesophagitis phenotypes; Outcome Ascertainment.
Covariates / confounders	Covariate inventory and rationale; concept-ID tables.
Missing data	Data-engineering / QC section.
Statistical methods	Logistic model, per-exposure full/backward AIC, AUC, VIF.
Multiplicity	Benjamini–Hochberg FDR q -values.
Sensitivity analyses	Distinct-date, temporal, overlap, treated-GERD.
Bias / validity threats	Threats to Validity and Limitations.
Reproducibility	Execution order; shared engine; software environment.
Data availability	Confidential; remains in the secure workspace.

Table 25: STROBE-aligned reporting and reproducibility checklist.

End of report. This document describes code and methodology only; it contains no individual-level *All of Us* data.